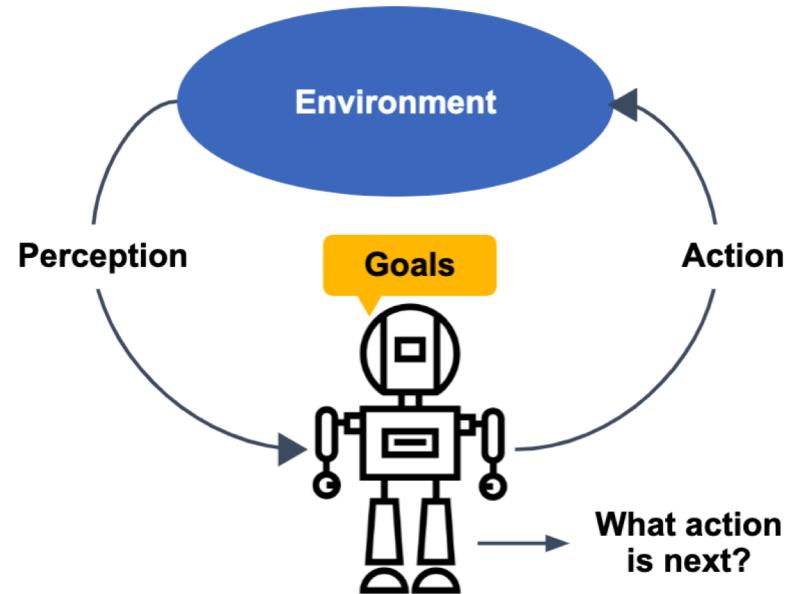

Sequential Decision-Making

Mehrdad Zakershahrak
Guest Speaker
Arizona State University

Sequential Decision Making

- | Intelligent agency involves controlling the evolution of external environments in desirable ways.
- | Planning provides a way in which the agent can maximize its chances of effecting this control.
- | Plan is course of actions that the agents decides upon based on its overall goals, information about the current state of the environment, and the dynamics of its evolution.



Classical Planning Problem

- | The complexity of **plan synthesis** depends on a variety of **properties of the environment** and the agent.
- | Perhaps the simplest case of planning occurs when the environment is **static**, (in that it changes only in response to the agents actions), **observable** and the agent's actions have **deterministic** effects on the state of the environment.
- | Plan synthesis under these conditions has come to be known as the classical planning problem.

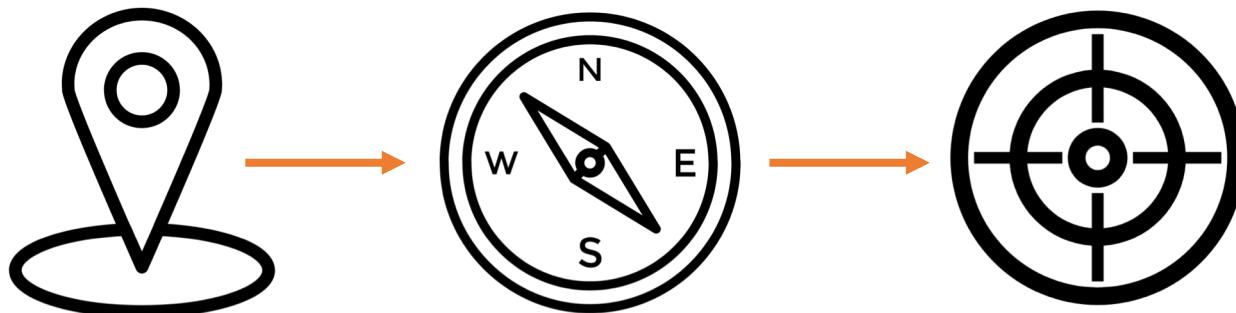
Classical Planning Problem

- | Classical planning problem is still **computationally hard**, and has received a significant amount of attention.
- | Work on classical planning has historically helped our understanding of planning under non-classical assumptions.

Example

| Pathfinding

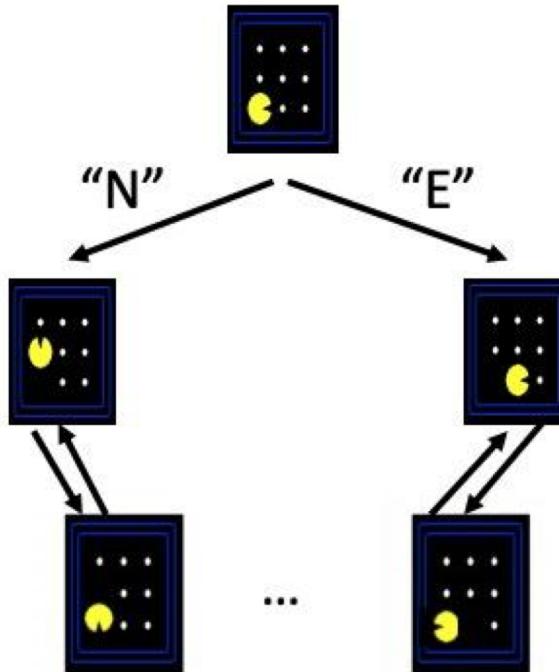
- States: (x,y) location
- Actions: NSEW
- Successor: updated location
- Goal: is (x,y) = END?



Search Problem

| State space graph: A **mathematical representation** of search problem

- Each node are *abstracted* environment configurations
- Edges are successors

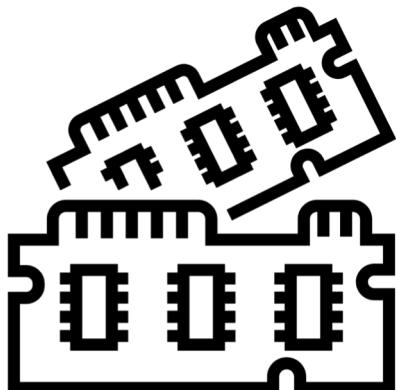


Search Problem

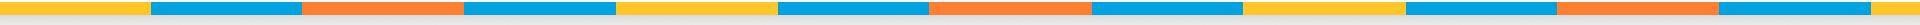
| Search tree

- a “What-if” tree of plans and their outcomes
- initial state is the root
- nodes correspond to states in the plan
- children correspond to successors

| Building the whole search tree takes lots of memory!



Search Strategies



| Order of **node expansion** matters

| Completeness

- Does it always find a solution if one exists?

| Optimality

- Does it always find a least-cost solution?

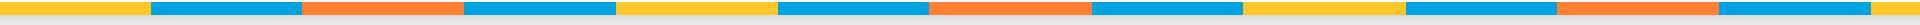
| Time complexity

- Number of nodes generated/expanded space

| Space complexity

- Maximum number of nodes in memory

Time and Space Complexity



| **Time and space complexity are measured with:**

- maximum branching factor
- depth of the shallowest solution
- maximum depth of the search space

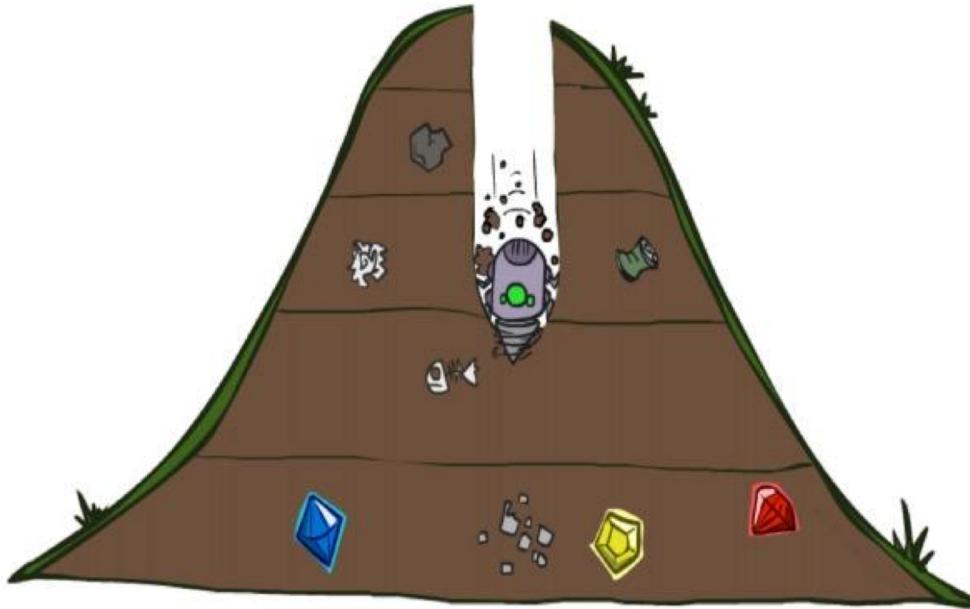
Uninformed Search

| Uninformed strategies use only the **information available** in the problem definition

- Depth-first search
- Breadth-first search
- Uniform-cost search
- Depth-limited search
- Iterative deepening search

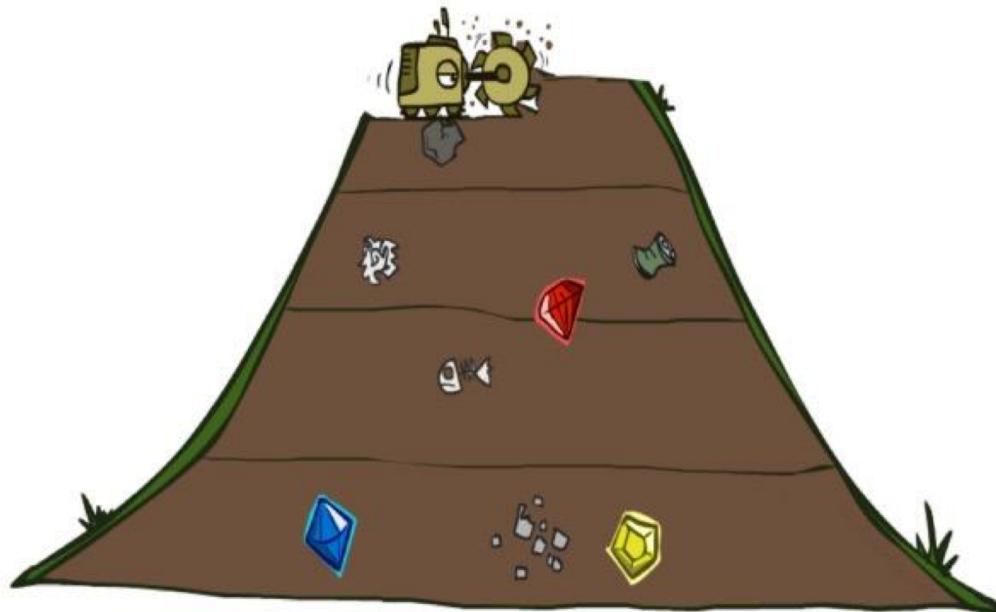
Depth-First Search

- | Strategy: Expand the deepest node first until going deeper is not possible. If the goal is not found, expand another node.
- | Implementation: Fringe is LIFO (Last In First Out) stack



Breadth-First Search

- | Strategy: expand a shallowest node first
- | Implementation: Fringe is a FIFO (First in First Out) queue



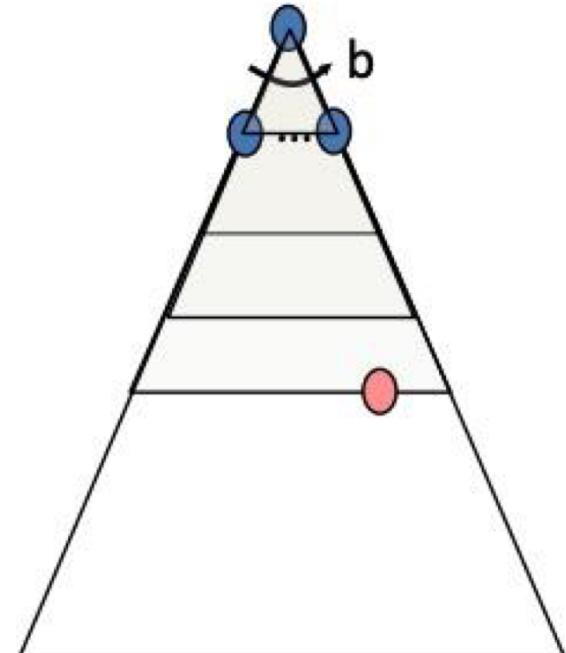
DFS vs. BFS

- | When will DFS outperform BFS?
- | When will BFS outperform DFS?

Iterative Deepening

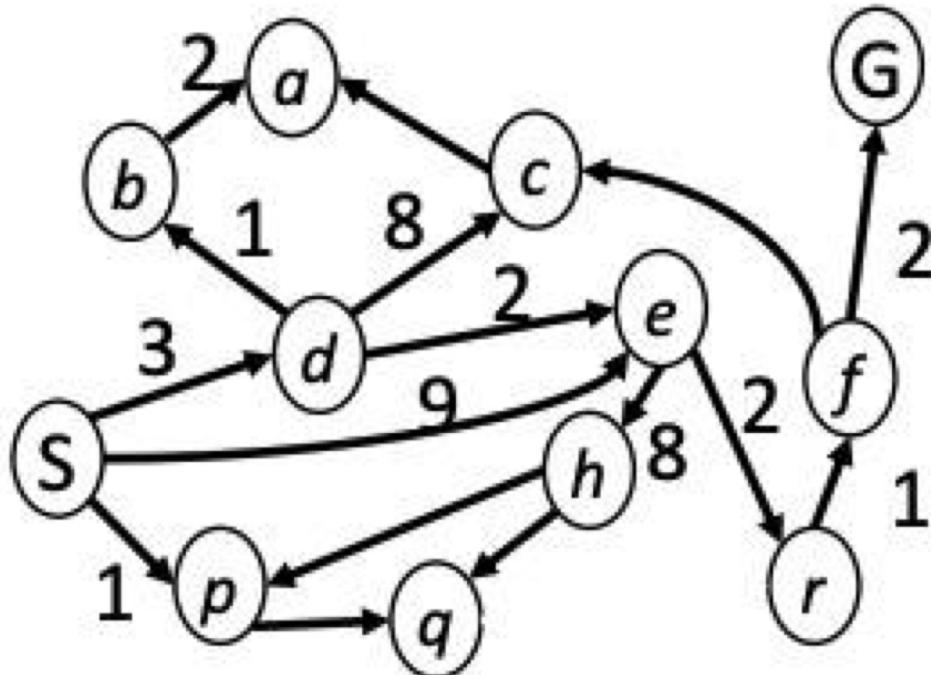
Idea: Get **DFS's space advantage** with **BFS's time / shallow-solution advantages**

- Run a DFS with depth limit 1. If no solution...
- Run a DFS with depth limit 2. If no solution...
- Run a DFS with depth limit 3



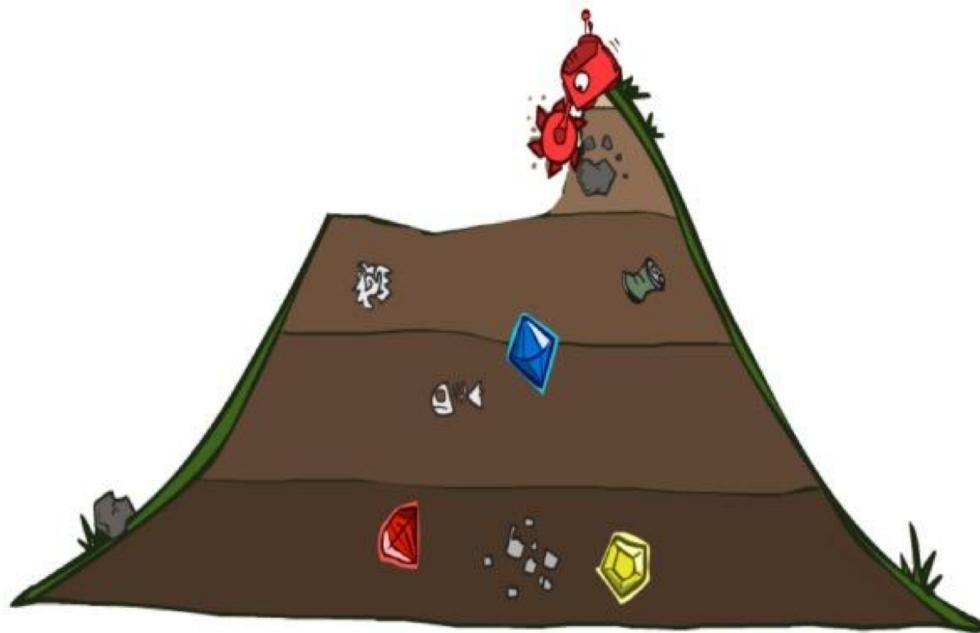
Cost Sensitive Search

- | BFS finds the shortest path in terms of number of actions.
- | It does not find the least-cost path.



Uniform Cost Search

- | Strategy: Expand the cheapest node first
- | Fringe is a priority queue (priority: cumulative cost)



Uniform Cost Search

- | UCS is **complete** and **optimal**
- | Explores options in every direction
- | No information about goal location

Summary

Criterion	Breadth-First	Uniform-Cost	Depth-First	Depth-Limited	Iterative Deepening
Complete?	Yes*	Yes*	No	Yes, if $l \geq d$	Yes
Time	$O(b^d)$	$O(b^{C^*/\varepsilon})$	$O(b^m)$	$O(b^l)$	$O(b^d)$
Space	$O(b^d)$	$O(b^{C^*/\varepsilon})$	$O(bm)$	$O(bl)$	$O(bd)$
Optimal?	Yes*	Yes	No	No	Yes*